

Pre-seed · 2026

CV COVER

Cover letters written from the real job posting,
not from a guess.

coveritt.vercel.app · live, deployed, no account required

Built on Bright Data · Into the Scrape-Verse (WeMakeDevs x Bright Data)

Applying to ten jobs means writing ten cover letters.

So people paste the ad into ChatGPT. That produces a letter that reads like ChatGPT wrote it: the same three-paragraph shape, the same "I am writing to express my enthusiasm", and an increasingly detectable statistical fingerprint that applicant tracking systems now screen for.

10x

LETTERS PER SEARCH

Every posting wants a letter that cites *that* posting. The work scales linearly with applications.

Generic

WHAT THE SHORTCUT RETURNS

A chatbot only sees what you pasted. Miss a detail from the ad, and the letter is filler.

Flagged

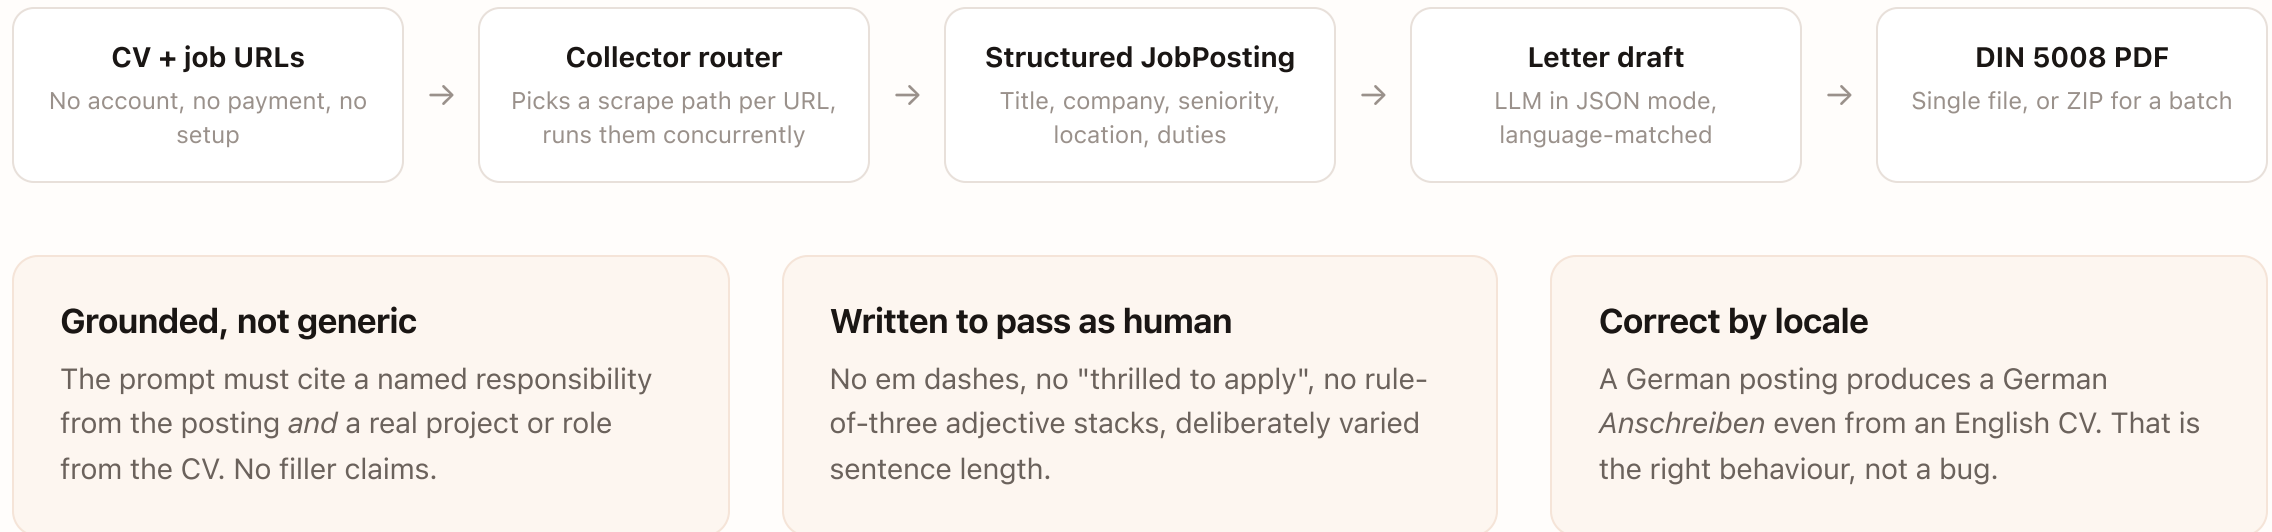
THE NEW FAILURE MODE

ATS-side AI detection turns the shortcut into a liability, not just a weak letter.

But the writing is the easy half. The hard half is acquiring the job posting reliably, as structured data, from a page built for human eyes and redesigned without notice. That acquisition problem is what this company is built around.

Upload a CV. Paste job links. Get one letter per posting.

Each letter is drafted from the actual scraped job ad and your real resume, in the posting's own language, and downloads as a formal PDF laid out to the German DIN 5008 letter standard.




Three inputs, one screen.

Drag in a CV, paste the job links, hit generate. There is no onboarding, no signup wall and no pricing page in the way. The entire product is one page and a results modal.

- **Any LinkedIn URL shape works** — including the messy one from your address bar with the job id buried in a query string.
- **Several jobs at once** — scraped concurrently, bundled into a ZIP.
- **Letterhead auto-fills** — name, address, phone, email and LinkedIn are parsed out of the CV itself.

Live at coveritt.vercel.app — screenshot captured from production.




Every scrape reports what it actually recovered.

Field by field, per run, in production. Note the status on these runs: **partial** — not a green checkmark. The postings listed no salary, so the run says so.

```
// real, unedited production record
{
  "collector_name": "linkedin_job",
  "status": "partial",
  "fields_recovered": ["role_title", "company_name",
    "seniority_level", "location", "responsibilities"],
  "fields_missing": ["salary"],
  "self_heal_events": []
}
```

A scraper that quietly degrades is worse than one that crashes, because the output still looks plausible. Naming the missing fields is how you tell the difference.



A scraper that fails loudly is a bug report. A scraper that fails *plausibly* sends your application to the wrong company.

What actually happened

The Scraper Studio collector was built against Greenhouse. Pointed at Lever and Ashby — unfamiliar layouts — it did **not** return empty fields.

It returned **the training company's name**, every time, with the record otherwise perfectly well-formed. The app happily generated letters addressed to the wrong employer.

The fix — at the application layer

Self-healing does not catch this, because from the collector's point of view nothing is broken. It found a company name.

```
# collectors/router.py
if not posting.role_title:
    raise BrightDataError(
        "Could not read this job posting.")
```

Treat a missing role title as proof the scrape did not land. Refuse rather than guess.

This is the moat in miniature. Anyone can call an LLM. The durable asset is a validated acquisition layer that knows when it is wrong — and a body of hard-won knowledge about how job boards fail.

Three collection paths, on purpose. One schema.

The router picks a path per URL. Everything downstream is path-agnostic, so adding a board is additive, never a rewrite.

URL	PATH	WHY THIS PATH
linkedin.com	Prebuilt LinkedIn Jobs dataset	LinkedIn is aggressively hostile to generic scraping and requires a session for most postings. The prebuilt dataset returns a fixed schema synchronously and sidesteps that entirely.
join.com	schema.org JobPosting collector	join.com renders client-side, so a generic collector reads nothing. Every posting ships a structured JSON-LD block — exact, and stable across layout changes.
everything else	Bright Data Scraper Studio collector	Fields are defined in plain language, not CSS selectors, so the collector can be repaired in place when a site's layout shifts.

FRONTEND

Next.js 16, Tailwind v4, client-side PDF + ZIP

BACKEND

FastAPI, httpx, pydantic, concurrent scrapes

ACQUISITION

Bright Data Scraper Studio + prebuilt datasets

DEPLOY

Vercel + Railway, live in production today

Three layers, cheapest first.

Scraping is not a one-time build. It is a maintenance liability, and the architecture is designed around that fact.

01

Plain-language fields

Extraction is driven by what a field *means* ("the job title for this posting"), not by a selector. A renamed CSS class does not silently produce empty strings.

02

Automated healing

`bdata scraper heal` repairs the collector against the live page when extraction genuinely regresses — no hand-written selectors to rewrite.

03

Guard against confident wrongness

The one that matters. Validate that what came back is actually about the page you asked for — and refuse when it is not.

Three run statuses, reported honestly

success every expected field came back

partial the scrape landed, some fields were absent from the page

failed nothing usable came back

What we refuse to do

Unsupported boards are rejected with a clear message rather than silently mis-scraped. Company context is optional and off by default, because it is unreliable on JS-heavy marketing sites. When it returns noise, the letter is written from the posting and CV alone rather than absorbing nonsense.

Bottom-up, from job applications actually written.

MODELLED — ASSUMPTIONS STATED

Sized from applicant behaviour rather than from a top-down "HR tech is \$30B" number, so every input below can be argued with individually.

TAM

\$2.4B

~**400M** white-collar job seekers globally ×
~**\$6/yr** blended willingness-to-pay for application tooling.

SAM

\$310M

~**26M** active seekers in EU + UK + North America who apply through the supported boards, at ~**\$12/yr** realised.

SOM — 3 YR

\$4.6M

1.5% of SAM. ~32k paying users at **\$144/yr** (see pricing), reached via the free tier already live.

Why the beachhead is German-language applications

DIN 5008 is a real barrier. German applications have a formal layout standard. Generic tools output an American business letter, which reads as wrong to a German recruiter.

Language matching is non-trivial. A German posting needs a German *Anschreiben* with the right conventions, not a translation. Already shipped.

join.com is a DACH board. Supporting it is a deliberate wedge into a market the US-built incumbents ignore.

Cost per letter, from the real call graph.

MODELLED — LIST PRICES

One letter = one scrape + one LLM draft. Both are metered, both are small, and the PDF is rendered client-side at zero marginal cost.

COST COMPONENT	PER LETTER
Bright Data scrape 1 collector run or dataset call	\$0.0060
LLM draft ~4k in / ~700 out, JSON mode	\$0.0100
CV parse + PDF render server CPU + client-side jsPDF	\$0.0004
Marginal COGS	\$0.0164

Scrape and draft costs are list-price estimates for the current call pattern, not negotiated rates or measured spend.

UNIT MODEL — PRO TIER	VALUE
Price	\$12 / mo
Letters included	40 / mo
Assumed real usage	14 / mo
COGS at that usage	\$0.23
Payment + infra overhead	\$0.72
Gross margin	92%

CAC

\$18

content + board SEO

LIFETIME

5 mo

job search is seasonal

LTV / CAC

3.1x

\$55 LTV

Free stays free. The batch is what converts.

PROPOSED — NOT YET LIVE

The product is free and account-less today. That is the top of the funnel, not the business model — and it is why the free tier keeps a real, unlimited-feeling allowance.

FREE

\$0

3 letters / month. No account. Full DIN 5008 PDF, no watermark, no degraded output.

Purpose: prove the letter quality is different before asking for anything.

PRO — THE CORE

\$12/mo

40 letters, ZIP batching, saved CV profiles, letter history.

The trigger is volume: the day someone applies to eight roles at once, the ZIP is worth \$12.

CAREERS SERVICES

\$450/yr

Per seat, university and bootcamp career offices. Cohort dashboards, institutional letterhead.

Solves the churn problem: the institution renews annually even as students cycle through.

Why not charge per letter? Per-letter pricing makes users ration the product at exactly the moment it is most useful — a big application push. Subscription pricing aligns with how job searches actually happen: nothing for months, then everything in three weeks.

Everyone has the LLM. Almost nobody has the acquisition layer.

The differentiator is not letter quality in the abstract. It is that the letter is written from a verified, structured posting instead of from whatever the user managed to paste.

APPROACH	GETS THE POSTING HOW	FAILS HOW	WHERE COVERIT DIFFERS
ChatGPT / Claude direct	User copy-pastes	Silently incomplete input	We scrape the full posting as structured fields, so nothing depends on what the user remembered to paste.
Cover-letter SaaS Teal, Kickresume, et al.	Paste, or a thin URL fetch	Generic US business-letter output	DIN 5008 layout and true language matching — a German posting yields a German Anschreiben.
LinkedIn Easy Apply	Native, no letter	No differentiation at all	We compete on the letter being specific, which is the only part a recruiter reads twice.
Build-it-yourself	Own scraper	Confidently wrong data	The failure mode on slide 06 is invisible until it has already cost someone an application. We built the guard.

Acquisition breadth

Three collection paths converging on one schema. Adding a board is additive.

Failure literacy

Per-field run records in production. We know when we are wrong, and we say so.

Locale correctness

DIN 5008 and language matching are shipped, not roadmap.

Built and deployed, not a prototype.

Everything below is live in production today. There is no paid tier and no user base yet — that is the honest position, and the next section is what changes it.

Shipped

- **Three collection paths** — LinkedIn (any URL shape), join.com, Greenhouse
- **Resume parsing** — PDF, TXT, Markdown, plus sender-block extraction for the letterhead
- **Language-matched generation** — one letter per job URL, in the posting's language
- **DIN 5008 PDF output** — individual files and ZIP bundling
- **Editable letters and letterhead** — review and correct before download
- **Collector health tracking** — per-field recovery, publicly visible
- **The confident-wrongness guard** — refuses rather than misattributes

Known limitations — stated, not hidden

- **Three boards supported.** Others are refused with a clear message rather than silently mis-scraped.
- **Extraction completeness varies** by layout. Incomplete runs report **partial** with fields named.
- **Contact extraction is regex-based.** Unusual CV layouts may miss a field; missing fields collapse rather than print blank.
- **Run history is in-memory** and resets on backend restart.
- **No paid tier, no auth, no user accounts** yet.


| The health view exposes these anyway. Stating them here is cheaper than being found out.

THE ASK

\$400k pre-seed — 18 months to a proven paid funnel.

PROPOSED

The engineering risk is retired: it works, it is deployed, and the hard failure mode is understood and guarded. What is unproven is willingness to pay. That is what this round buys an answer to.



MILESTONE	TARGET	WHAT IT PROVES
Month 3	Accounts, billing, 15 job boards	The acquisition layer generalises beyond the three boards it was built against.
Month 6	1,000 paying users	The \$12 price point clears against a genuinely useful free tier.
Month 12	First 10 careers-services contracts	The institutional channel solves structural churn.
Month 18	\$45k MRR, LTV/CAC ≥ 3	The unit model on slide 10 survives contact with real users.

The writing was never the hard part. Knowing the posting is *real* is.

Coverit is a validated job-posting acquisition layer with a cover-letter product on top. The letter is what users buy. The acquisition layer is what competitors cannot copy from a screenshot.

LIVE PRODUCT

coveritt.vercel.app

COLLECTOR HEALTH

</health>

SOURCE

github.com/karankartikeya/cvsmart

Built for Into the Scrape-Verse · WeMakeDevs × Bright Data